



មហាវិទ្យាល័យវិស្វកម្ម
FACULTY OF ENGINEERING

Introduction to Deep Learning Applications and Theory

Lecture 5 ANN Implementation in Python

Chhoeum Vantha, Ph.D.
Telecom & Electronic Engineering

Previous week: ANN's Theory


- Introduction to artificial neural network (ANN)
- The Perceptron

Outline

- Layers: Python Operators
- Optimizer
- Summarize Model
- Training Loss
- Validation Loss
- Designing Training Procedures


Layers: Python Operators

- Activation Functions
- Normalization
- Dropout
- Loss Functions




Layers: Python Operators

- Activation Functions
- Normalization
- Dropout
- Loss Functions




Activation Functions

- Non-linear functions performed by neurons
- ReLU - Rectified Linear Unit (nn.ReLU)
 - $y \geq 0$
- Tanh (nn.tanh)
 - $-1 < y < 1$
 - nn.Tanh
- Sigmoid (nn.Sigmoid)
 - $0 < y < 1$



Activation Functions


- Leaky ReLU
 - Similar to ReLU, but has non-zero values for negative x
 - Takes argument *negative_slope*, which determines the slope for $x < 0$.
- For full list of activation functions, see: <https://pytorch.org/docs/stable/nn.html>



Leaky ReLU with negative slope = 0.1


Layers: Python Operators

- Activation Functions
- **Normalization**
- Dropout
- Loss Functions



Input Normalization: Batch Normalization


- **Normalizes input** into each layer for each training mini-batch
- Addresses issue of **shifting input** distributions over training
- **Batch normalization** applies a transformation that maintains the mean output close to **0** and the output standard deviation close to **1**



Input Normalization: Batch Normalization

- Batch Normalization

```
tf.keras.layers.BatchNormalization(  
    axis=-1,  
    momentum=0.99,  
    epsilon=0.001,  
    center=True,  
    scale=True,  
    beta_initializer='zeros',  
    gamma_initializer='ones',  
    moving_mean_initializer='zeros',  
    moving_variance_initializer='ones',  
    beta_regularizer=None,  
    gamma_regularizer=None,  
    beta_constraint=None,  
    gamma_constraint=None,  
    synchronized=False,  
    **kwargs  
)
```




Input Normalization: Batch Normalization

- **Inputs:**
 - **num_features:** Number of features in the input vector
 - **eps:** numerical stability parameter

```
from keras.models import Sequential
from keras.layers import Dense, BatchNormalization

model = Sequential()
model.add(Dense(32, input_shape=(64,)))
#This is the Batch normalization layer
model.add(BatchNormalization())
model.add(Dense(32))
model.add(BatchNormalization())
model.add(Dense(10, activation='softmax'))
```




Input Normalization

Other normalization procedures include:

- **Layer Norm**: Transposes Batch Norm.
- Normalizes over all summed inputs to a layer
 - <https://arxiv.org/abs/1607.06450>
- **Group Norm**: Normalizes by grouped channels instead of batches
 - <https://arxiv.org/abs/1803.08494>


Layers: Python Operators

- Activation Functions
- Normalization
- Dropout
- Loss Functions




Dropout

- The term “dropout” refers to **dropping out the nodes** (input and hidden layer) in a neural network.
- All the forward and backwards connections with a dropped node are **temporarily removed**, thus creating a new network architecture out of the parent network.
- **The nodes are dropped** by a dropout probability of p .




(a) Standard Neural Net



(b) After applying dropout.


Dropout

- **Randomly zeroes** some elements of input tensor with probability p
- Effective technique for regularization Outputs scaled by $1/1-p$
- Treated as identity during evaluation




Layers: Python Operators

- Activation Functions
- Normalization
- Dropout
- Loss Functions



Loss Functions

- Loss function parameters:
 - **Reduction**: how the output will be reduced in dimension:
 - **None**: Gives entire Loss Tensor with no reduction over batches
 - **Sum**: Takes sum of the loss tensor across batches, returning a single number
 - **Mean**: Same as sum, but divides by the number of batches to get the mean



Loss Functions: keras

Probabilistic losses

- BinaryCrossentropy class
- CategoricalCrossentropy class
- SparseCategoricalCrossentropy class
- Poisson class
- binary_crossentropy function
- categorical_crossentropy function
- sparse_categorical_crossentropy function
- poisson function
- KLDivergence class
- kl_divergence function

Regression losses

- MeanSquaredError class
- MeanAbsoluteError class
- MeanAbsolutePercentageError class
- MeanSquaredLogarithmicError class
- CosineSimilarity class
- mean_squared_error function
- mean_absolute_error function
- mean_absolute_percentage_error function
- mean_squared_logarithmic_error function
- cosine_similarity function
- Huber class
- huber function
- LogCosh class
- log_cosh function

<https://keras.io/api/losses/>

Loss Functions – Cross Entropy Loss

- Cross Entropy Parameters



- **Weight**

- 1D tensor assigning weights to each class, which is helpful if you have an unbalanced training set

- **ignore_index**

- Specifies a target value that is ignored and does not contribute to the input gradient

```
keras_core.losses.BinaryCrossentropy(  
    from_logits=False,  
    label_smoothing=0.0,  
    axis=-1,  
    reduction="sum_over_batch_size",  
    name="binary_crossentropy",  
)
```




Optimizer

- **Optimizers** are algorithms or methods used to minimize an error function(loss function)or to maximize the efficiency of production.
- Optimizers are mathematical functions which are dependent on model's learnable parameters i.e **Weights & Biases**.
- Optimizers help to know how to change weights and learning rate of neural network to reduce the losses.


Optimizer: Initialization

- **Parameters:**
 - Should be iterable containing parameters to optimize
 - E.g., `model.parameters()` or `[var1, var2]`
 - **Parameters** must be defined BEFORE the optimizer
- **Optimizer Settings**
 - Learning rate, weight decay, etc.




Optimizer: Initialization-Param

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 250)	25250
dense_1 (Dense)	(None, 150)	37650
dense_2 (Dense)	(None, 250)	37750
dense_3 (Dense)	(None, 100)	25100
Total params: 125750 (491.21 KB)		
Trainable params: 125750 (491.21 KB)		
Non-trainable params: 0 (0.00 Byte)		



Optimizer – Functions in Keras

- SGD
- RMSprop
- Adam
- AdamW
- Adadelta
- Adagrad
- Adamax
- Adafactor
- Nadam
- Ftrl




<https://keras.io/api/optimizers/>


Optimizers

- Stochastic Gradient Descent
(`optim.SGD`)
 - **params**: Model parameters
 - **lr**: Learning rate (required)
 - **momentum**: momentum factor (default: 0)
 - **weight_decay**: (default: 0)

SGD without momentum



SGD with momentum



Optimizers

- Adam ()
 - **params**: Model parameters
 - lr: Learning rate (default: 0.001)
 - betas: coefficients (tuple) used for computing running averages of gradient and its square (default: (0.9, 0.999))
 - eps: term added to denominator to improve numerical stability (default: 1e-8)
 - weight_decay: (default: 0)

Other Common Optimizers

- AdaDelta()
 - Precursor to Adam which uses first-order estimates to adapt learning rate
- Adamax()
 - Variant on Adam based on infinity norm
- RMSProp()
 - Take the square root of the gradient average before adding epsilon to normalization of LR

```
# Add Optimizer to minimize the loss
model1.compile(loss='mean_squared_error', optimizer=RMSprop(lr=0.007), metrics=['accuracy'])
```

Summarize Model

- Keras provides a way to summarize a model.
- The summary is textual and includes information about:
 - The layers and their order in the model.
 - The output shape of each layer.
 - The number of parameters (weights) in each layer.
 - The total number of parameters (weights) in the model.

Training Loss

- The training loss is a metric used to assess how a deep learning **model** **fits the training data**
- It assesses the error of the model on the training set.
- **The training loss** is calculated by taking the **sum of errors** for each example in the training set.

Validation Loss

- Validation loss is a metric used to assess the **performance** of a deep learning **model** on the **validation set**.
- The validation loss is similar to the training loss and is calculated from a sum of the errors for each example in the validation set.

Practice 5

Deep Neural Network procedure in Python code:

```
# 1. Import the library
# 2. Load data
    # 2.1 Input data
    # 2.2 Target Data
# 3. Pre-processing data
    # 3.1 Convert data to array
    # 3.2 Transpose the matrix
# 4. Normalize the data
    # 4.1 formula to normalize
```

Deep Neural Network procedure in Python code:

5. Label the target output

 # 5.1 formula to label the target output

6. Split data 80% for train and 20% for validation

7. Create deep learning model

8. Compile and train model

9. Visualize loss in deep learning

10. Visualize accuracy in deep learning

11. Save the models

11. Test model

12. Check the performance of model with actual output and predict output

Thanks!